

Cloud Computing - Project Report

Totoro Group

Zahra Omrani
z.omrani@studenti.unipi.it

Paolo Palumbo
p.palumbo3@studenti.unipi.it

Ettore Ricci
e.ricci32@studenti.unipi.it

June 15, 2024

<https://github.com/Etto48/CloudComputingProject>

Abstract

This report presents the development and performance evaluation of a Java-based application designed to count letter frequencies in large datasets using the Hadoop framework. To provide a comprehensive analysis, we also implemented and tested non-distributed applications in Python and Rust for comparison. The results show that for the sizes of the datasets considered, Java takes less time to complete than Python, while Rust is significantly faster than Java. Both Python and Rust applications used significantly less memory than the Java application.

1 Introduction

Our project for the Cloud Computing course consists of developing a Java application to count the frequency of letters in a large dataset using the Hadoop framework. The main objective of the project is to analyze and compare the performance of the application with different configurations and input sizes and also to compare it with non-distributed implementations in Python and Rust.

- **Mapper with in-mapper combiner:** This mapper stores the letter counts in a vector and emits them in the cleanup method. At the end of the cleanup method, the mapper increments a job counter with the number of emitted letters.
- **Mapper without in-mapper combiner:** This mapper emits the letter counts (with value 1) for each letter in the map method. At the end of the map method, the mapper increments a job counter with the number of emitted letters.

2 Mapreduce

MapReduce is a programming model for processing large data sets. Users specify a map function that processes a key/value pair to generate a set of intermediate key/value pairs, and a reduce function that merges all intermediate values associated with the same intermediate key. For our application we needed two subsequent MapReduce jobs: the first one to count the occurrences of each letter in the input text and the second one to calculate the frequency of each letter.

2.1 Job 1: Counting letters

2.1.1 Mapper

For testing purposes, we created two different mappers: one with an in-mapper combiner and one without. For the combiner (when enabled), we used the same code as the reducer.

2.1.2 Reducer

The reducer receives the letter counts from the mappers and sums them up.

2.2 Job 2: Calculating frequencies

The results of the first job are saved in a temporary file and used as input for the second job. The counter obtained from the first job is stored in a configuration variable that is passed to the reducer (only one is needed because the number of different letters is small). For this job we do not need a combiner so it's always disabled.

2.2.1 Mapper

We used an identity mapper that emits the letter counts as they are.

2.2.2 Reducer

The reducer receives the letter counts from the mappers and calculates the frequency of each letter using the total count that was previously stored in a configuration variable.

that reads the input file in an asynchronous way, discarding the parts of the file that have already been processed (file streaming), and one that first reads the whole file into memory and then processes it.

3 Dataset

We used three different datasets:

- **english.txt**: A 1.2GB text file containing text from the books on [Gutenberg project](#).
- **italian.txt**: A 1.3GB text file from the [PAISÀ corpus](#)[1].
- **spanish.txt**: A 130MB text file from the [Leipzig Corpora Collection](#)[2].

We also created 11 files of increasing size (from 100MB to 1.1GB with 100MB steps) to test the performance of the application with different input sizes. These files are called **part_XMB.txt** and contain the first XMBs of the **english.txt** file.

4 Experiments

We wrote our program in a way that allows us to configure it for the different tests with command line arguments:

- `-i` | `--input`: Input file path inside the HDFS.
- `-r` | `--reducers`: Number of reducers to use in the first job.
- `-n` | `--no-combiner`: Disable the combiner in the first job.
- `-m` | `--no-in-mapper-combiner`: Disable the in-mapper combiner in the first job.
- `-o` | `--output`: Output directory path inside the HDFS, by default it's set to **output**.
- `-t` | `--tmp`: Temporary directory path inside the HDFS, by default it's set to **tmp**.

We carried out 21 tests with the configurations shown in Table 1 and Table 2. Originally we ran tests 7, 8 and 9 with 4 reducers, but test 9 could never complete because of a Shuffle Error caused by a Java Heap Space error (Out Of Memory Error). Because of this, we decided to run these tests again with 3 reducers. Only the 3 reducers configuration will be shown in the results. We also ran the Python and Rust applications with the same input file (english.txt). We tested two slightly different versions of the Python and Rust applications: one

Test ID	Arguments
0	<code>-i english.txt</code> <code>-r 1</code>
1	<code>-i english.txt</code> <code>-r 2</code>
2	<code>-i english.txt</code> <code>-r 4</code>
3	<code>-i english.txt</code> <code>-r 8</code>
4	<code>-i english.txt</code> <code>-r 1</code> <code>--no-combiner</code>
5	<code>-i english.txt</code> <code>-r 1</code> <code>--no-in-mapper-combiner</code>
6	<code>-i english.txt</code> <code>-r 1</code> <code>--no-in-mapper-combiner</code> <code>--no-combiner</code>
7	<code>-i english.txt</code> <code>-r 3</code> <code>--no-combiner</code>
8	<code>-i english.txt</code> <code>-r 3</code> <code>--no-in-mapper-combiner</code>
9	<code>-i english.txt</code> <code>-r 3</code> <code>--no-in-mapper-combiner</code> <code>--no-combiner</code>

Table 1: General tests

Test ID	Arguments
10	-i part_100MB.txt -r 1
11	-i part_200MB.txt -r 1
12	-i part_300MB.txt -r 1
13	-i part_400MB.txt -r 1
14	-i part_500MB.txt -r 1
15	-i part_600MB.txt -r 1
16	-i part_700MB.txt -r 1
17	-i part_800MB.txt -r 1
18	-i part_900MB.txt -r 1
19	-i part_1000MB.txt -r 1
20	-i part_1100MB.txt -r 1

Table 2: Input split tests

4.1 Equipment

The tests were carried out on 3 identical virtual machines, each one with the following specifications:

- **CPU:** 1 core (2 logical cores) of an Intel(R) Xeon(R) Silver 4208 CPU @ 2.10GHz
- **RAM:** 6.8GB
- **Disk:** 40GB
- **OS:** Ubuntu 22.04.4 LTS

The Java application used the whole equipment, while the Python and Rust applications used only one logical core on one machine.

5 Results

The results of the tests on the Java application are shown in Figure 4. The memory usage over time of the Python and Rust applications is shown in Figure 5. The frequencies of the letters calculated by the application are shown in Figure 1 for English, Figure 2 for Italian and Figure 3 for Spanish.




Figure 1: English letter frequency




Figure 2: Italian letter frequency




Figure 3: Spanish letter frequency

5.1 Java

5.1.1 Execution time

From Figure 4 (Execution time) we can see that the execution time of the Java application with the

complete dataset as input (Tests [0, 9]) has a minimum value of 76s and a maximum value of 893s. Execution time increases significantly when the in-mapper combiner is not used (Tests 5, 6, 8, 9) and increases further when the combiners are also disabled (Tests 6, 9). More reducers lead to a slight increase in execution time when using the in-mapper combiner (Tests 0, 1, 2, 3). Execution time of Test 9 is smaller than the one of Test 6 because the higher number of reducers helps with the processing speed of the big amount of data produced by the mappers without in-mapper combiners nor combiners. In all other cases, the overhead of more reducers is not compensated by the speedup in processing. We suppose that a big factor in this overhead is the memory usage increase due to the necessity of a shuffle phase. On the other hand, we can see from Tests [10, 20] that the execution time is almost constant (56s) until the input size reaches 600MB (Test 15), after which it starts to increase linearly up to 76s for the 1.2GB input file (Test 0).

5.1.2 Input splits

From Figure 4 (Splits) we can see that the number of input splits increases linearly with the input size, with a minimum of 1 for the 100MB input file (Test 10) and a maximum of 9 for the 1.2GB input file (Test 0). The number of input splits is not affected by the number of reducers nor by the presence of the in-mapper combiner or the combiner. The number of input splits of the second job is linearly proportional to the number of reducers of the first job, in our case it is always equal to the number of reducers. This is true because the dimensions of the intermediate files is small.

5.1.3 Memory usage

We can see from the top two graphs in Figure 4 that the memory usage of the Java application is very

high, reaching 61.3GB of virtual memory in Test 3. The memory usage increases with the number of reducers. The absence of the in-mapper combiner and the combiner does not significantly affect the total virtual memory usage. On the other hand, it significantly increases the virtual memory usage of the reducers. We think that this is due to the fact the most of the memory used by the mappers is released before the reducers start. We also see that the total memory usage scales linearly with the input size: with a minimum of 7.2GB for the 100MB input file (Test 10) and a maximum of 36GB for the 1.2GB input file (Test 0). When we used 1 reducer, the memory usage of the second job was negligible compared to the first job. With more reducers, we had more input files for the second job, which increased the linearly the memory usage of the second job, up to 32.4GB for 8 reducers (Test 3).

5.2 Python

The Python application completed after 293s using file streaming and 298s without it, taking approximately 4 times more than the Java application (Test 0). The memory usage is of 0.27GB with file streaming and 2.65GB without it, using up to 133 times less memory than the Java application (Test 0).

5.3 Rust

The Rust application completed after 24s using file streaming and 19s without it, taking 3 to 4 times less than the Java application (Test 0). The memory usage is of 0.0075GB with file streaming and 4.26GB without it, using up to 4800 times less memory than the Java application (Test 0). These results show the efficiency of Rust in terms of both execution time and memory usage and the huge overhead of an interpreted language plus the Hadoop framework.

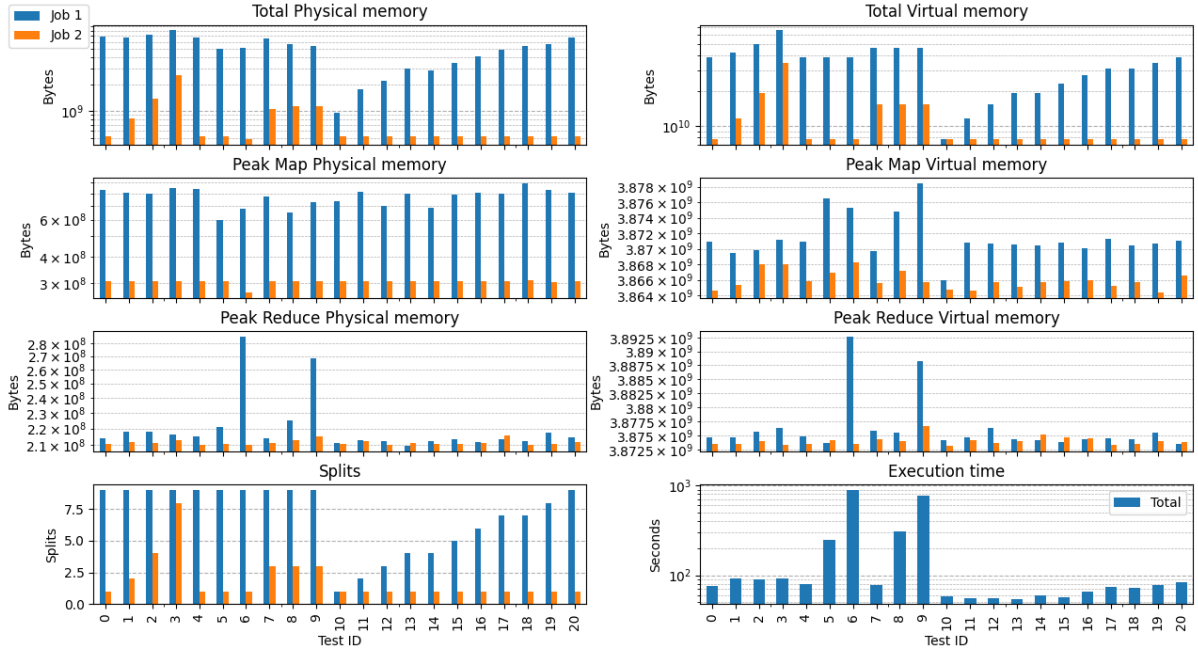


Figure 4: Tests results, the x-axis represents the test ID.

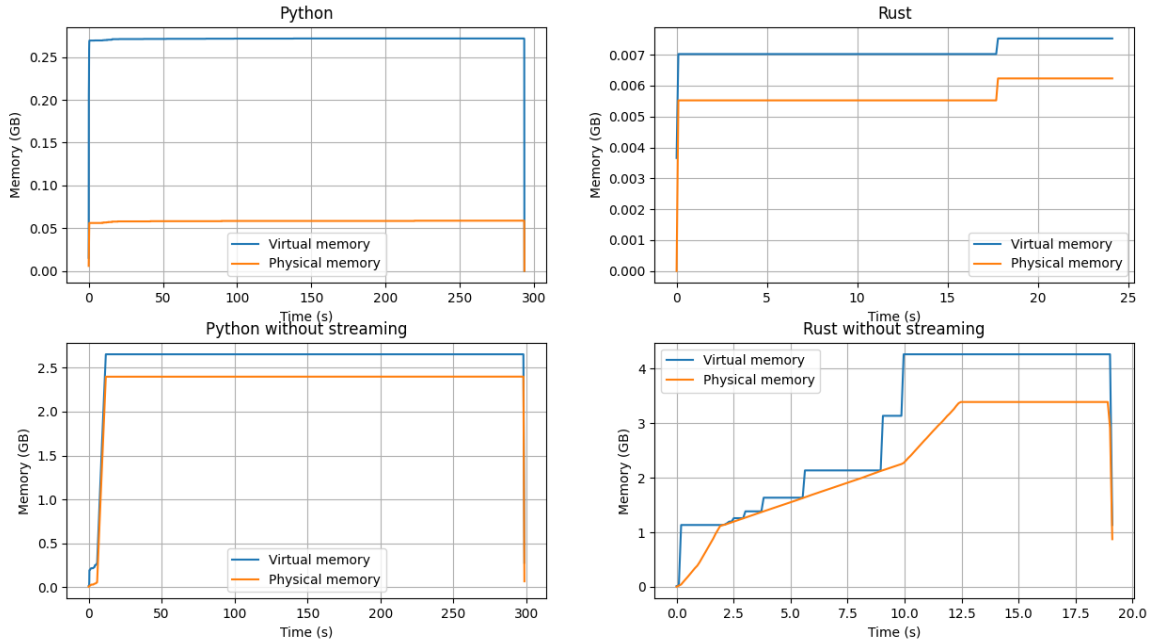


Figure 5: Python and Rust memory usage over time, with and without file streaming.

6 Conclusions

In this work we developed a Java application to count the frequency of letters in a large dataset using the Hadoop framework. We also developed non-distributed applications in Python and Rust for comparison. We carried out 21 tests with different configurations and input sizes one the Java

application and one test for each Python and Rust application then we analyzed and compared the results. The Java application is faster than the Python application and slower than the Rust application. The Rust application proved to be the most efficient in terms of both execution time and memory usage, using up to 4800 times less memory than the Java application. The Python application

still used significantly less memory than the Java application. Further improvements could be made to augment the precision of the results, such as performing more tests to have a more accurate average execution time and memory usage. The results that

a non-distributed, single-core, Rust application has obtained, suggest that a distributed Rust application could be significantly more efficient than an Hadoop application even for bigger datasets.

References

- [1] V. Lyding, E. Stemle, C. Borghetti, M. Brunello, S. Castagnoli, F. Dell’Orletta, H. Dittmann, A. Lenci, and V. Pirrelli, “The PAISÀ corpus of Italian web texts,” in *Proceedings of the 9th Web as Corpus Workshop (WaC-9)*, F. Bildhauer and R. Schäfer, Eds. Gothenburg, Sweden: Association for Computational Linguistics, Apr. 2014, pp. 36–43. [Online]. Available: <https://aclanthology.org/W14-0406>
- [2] T. Eckart and U. Quasthoff, *Statistical Corpus and Language Comparison on Comparable Corpora*. Berlin, Heidelberg: Springer Berlin Heidelberg, 2013, pp. 151–165. [Online]. Available: https://doi.org/10.1007/978-3-642-20128-8_8